

<b>Berkeley Socket API</b>	<b>2</b>
Byte Order	2
htons	2
htonl	2
ntohs	2
ntohl	2
<b>Addresses</b>	<b>3</b>
IP	3
IPv4	3
IPv6	4
Unix Domain	4
Address conversion	4
Interfaces	4
Functions	5
socket	5
bind	6
listen	6
connect	6
accept	6
read	6
write	7
close	7
recv	7
send	7

# Berkeley Socket API

## Byte Order

A host could use little-endian or big-endian.

The Berkeley Socket API uses network byte order, which is big-endian. Hosts must convert from host byte order (little-endian or big-endian) to network byte order (big-endian).

There are conversion functions provided to convert between host and network byte order.

If the host is big-endian, then these conversion functions do nothing.

Always call the conversion functions to ensure your code is portable, even if you are currently on a big-endian machine.

### htons

[`uint16\_t htons\(uint16\_t hostshort\)`](#)

Convert a 16-bit value from host byte order to network byte order.

### htonl

[`uint32\_t htonl\(uint32\_t hostlong\)`](#)

Convert a 32-bit value from host byte order to network byte order.

### ntohs

[`uint16\_t ntohs\(uint16\_t netshort\)`](#)

Convert a 16-bit value from network byte order to host byte order.

### ntohl

[`uint32\_t ntohl\(uint32\_t netlong\)`](#)

Convert a 32-bit value from network byte order to host byte order.

# Addresses

## IP

```
struct sockaddr
{
    sa_family_t ss_family
    char        sa_data[]
}
```

The `ss_family` is at least

Family	Purpose
AF_INET	IPv4
AF_INET6	IPv6
AF_UNIX	UNIX Domain socket, local to the machine

The `sa_data` is a memory trick in C. It makes the struct large enough to hold one of the structs below.

## IPv4

```
struct sockaddr_in
{
    sa_family_t    sin_family        // AF_INET
    in_port_t      sin_port          // network byte order
    struct in_addr sin_addr
}

struct in_addr
{
    uint32_t s_addr                // network byte order
}
```

## IPv6

```
struct sockaddr_in6
{
    sa_family_t    sin6_family    // AF_INET6
    in_port_t      sin6_port      // network byte order
    uint32_t       sin6_flowinfo
    struct in6_addr sin_addr
    uint32_t       sin6_scope_id
}

struct in6_addr
{
    uint8_t s6_addr[16]          // network byte order
}
```

## Unix Domain

```
struct sockaddr_un
{
    sa_family_t    sun_family    // AF_UNIX
    char           sun_path[]
}
```

## Address conversion

[`in\_addr\_t inet\_addr\(const char \*cp\)`](#)

Converts a string (e.g. "192.168.0.7") to a `inetaddr_t`.

## Interfaces

There is no way to get a list of network interfaces in POSIX. However, most Unix-like systems provide the `getifaddrs` function.

[`int getifaddrs\(struct ifaddrs \*\*ifap\)`](#)

The `ifap` argument is a pointer to a linked list that will be filled in. Each interface is an entry in the linked list. Interfaces can appear multiple times for each family they have.

A network interface can handle IPv4 and IPv6 simultaneously.

## Functions

Client	Server
<a href="#">socket</a>	<a href="#">socket</a>
	<a href="#">bind</a>
	<a href="#">listen</a>
<a href="#">connect</a>	<a href="#">accept</a>
<a href="#">read/write/recv/send</a>	<a href="#">read/write/recv/send</a>
<a href="#">close</a>	<a href="#">close</a>

## socket

[int socket\(int domain, int type, int protocol\)](#)

Create an unbound socket.

Both the client and server must create sockets.

A socket must be bound to a port.

The domains are, at least:

Domain	Purpose
AF_INET	IPv4
AF_INET6	IPv6
AF_UNIX	UNIX Domain socket, local to the machine

The types are, at least:

Type	Purpose
SOCK_STREAM	Stream socket
SOCK_DGRAM	Datagram socket
SOCK_RAW	Raw socket

The following combinations are valid:

	AF_INET	AF_INET6	AF_UNIX
<b>SOCK_STREAM</b>	TCP	TCP	Yes
<b>SOCK_DGRAM</b>	UDP	UDP	Yes
<b>SOCK_RAW</b>	IPv4	IPv6	

The protocol should be set to 0. Non-zero options are supported but are not needed for anything we will be doing (for now).

## bind

```
int bind\(int socket, const struct sockaddr \*address, socklen\_t address\_len\)
```

Assign a local address to a socket so a server can listen for client connections.

## listen

```
int listen\(int socket, int backlog\)
```

Used by a server to listen for clients initiating a 3-way handshake.

The backlog is the number of connections that can be queued.

## connect

```
int connect\(int socket, const struct sockaddr \*address, socklen\_t address\_len\)
```

Typically used to establish a TCP connection from a client to a server.  
Starts the 3-way handshake by sending the server a SYN packet.

## accept

```
int accept\(int socket, struct sockaddr \*restrict address, socklen\_t \*restrict address\_len\)
```

The server calls accept, which blocks, waiting for a client to connect to it. Accepts sends a SYN/ACK packet. Once the client receives the SYN/ACK it responds with a SYN packet and the 3-way handshake is complete.

## read

```
ssize\_t read\(int fildes, void \*buf, size\_t nbyte\)
```

The read function reads from a file descriptor. The accept and connection functions return a file descriptor where the file entry is a socket. Reading from the file descriptor reads data from the network connection.

The client and server can read in any order, based on the protocol.

## write

[ssize\\_t write\(int fildes, const void \\*buf, size\\_t nbytes\)](#)

The write function writes to a file descriptor. The accept and connection functions return a file descriptor where the file entry is a socket. Writing to the file descriptor writes data to the network connection.

The client and server can write in any order based on the protocol.

## close

[int close\(int fildes\)](#)

Close shuts down the network connection with the FIN and ACK sequence.

The client or the server can close first, based on the protocol.

## recv

[ssize\\_t recv\(int socket, void \\*buffer, size\\_t length, int flags\)](#)

The recv function is the same as read, but it only works with sockets, and you can pass the flags to provide some additional control.

## send

[ssize\\_t send\(int socket, const void \\*buffer, size\\_t length, int flags\)](#)

The send function is the same as write, but it only works with sockets, and you can pass the flags to provide some additional control.